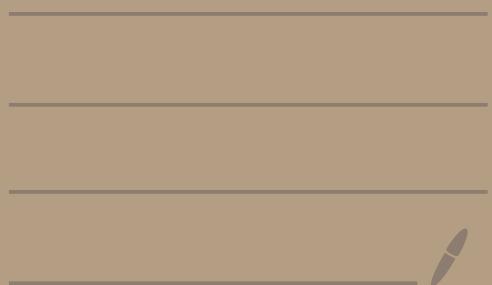


Computer Architecture Unit 2

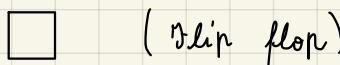


22/02/2022

BYTES

(INTRODUCTION)

bit



byte (8 bits)



kilobyte : 10^3 bytes

Kilobyte 1024 bytes

Megabyte : 10^6 bytes

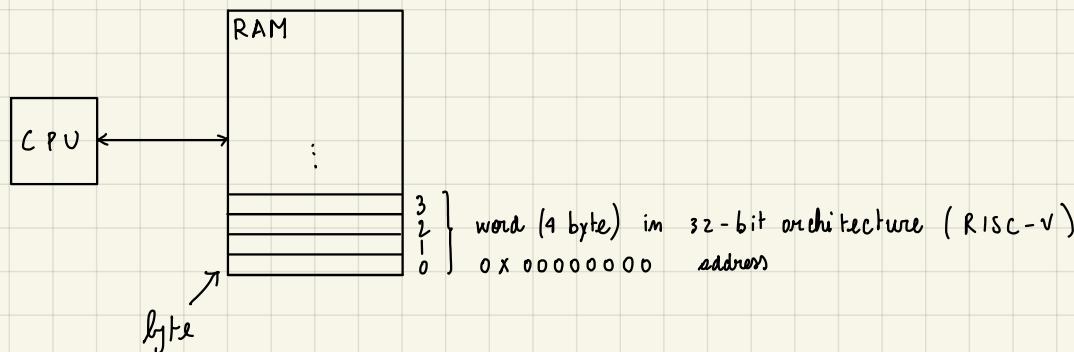
Megabyte 1024 kilobytes

Gigabyte : 10^9 bytes

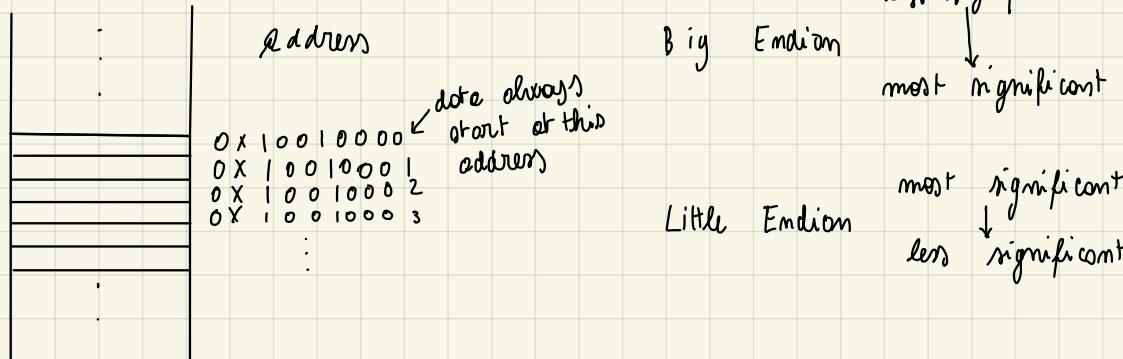
Gigabyte 1024 Megabytes

Terabyte : 10^{12} bytes

Terabyte 1024 Gigabytes

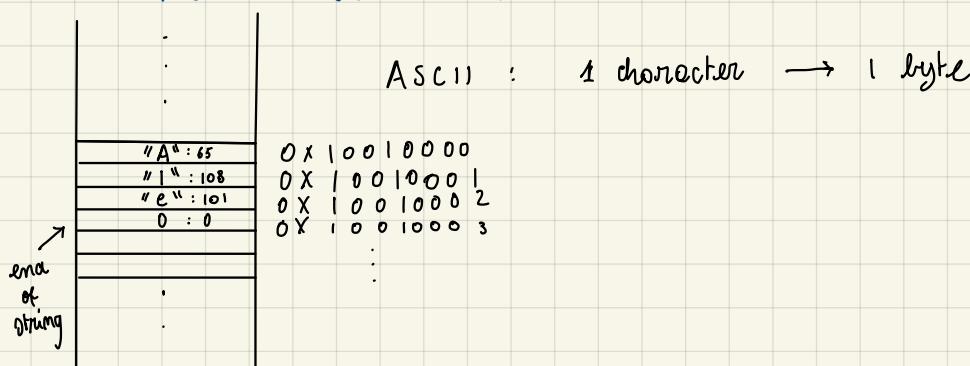


STORING INTEGERS IN RAM



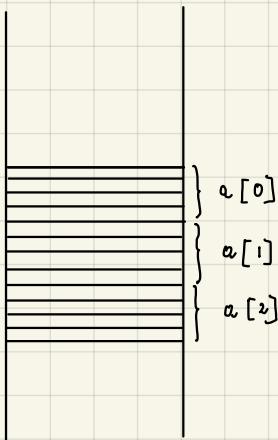
- 2-complement is used
- 4-bytes $\rightarrow 2^{32}$ numbers $\rightarrow [-2^{32}, +2^{32}-1]$

STORING STRINGS IN RAM



STORING ARRAYS IN RAM

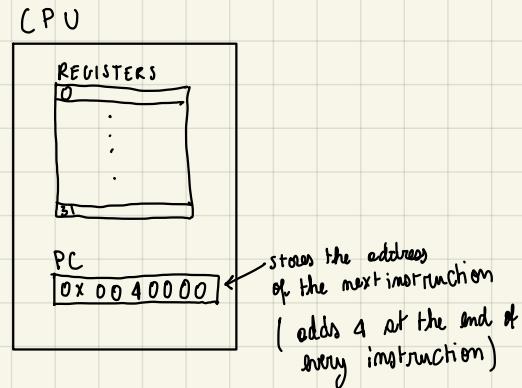
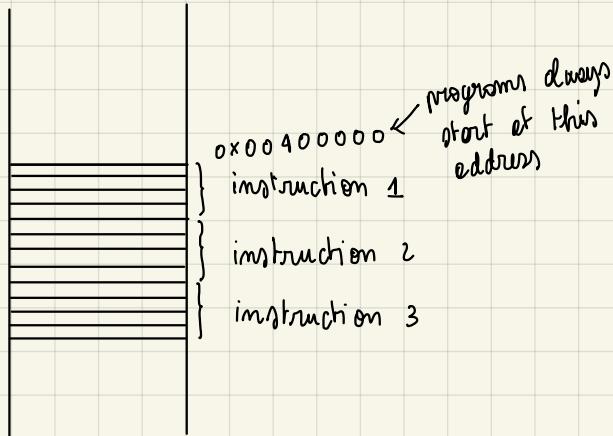
$$a = [7, 2, -3]$$



- we have to store also the lenght of the array

STORING PROGRAMS IN RAM

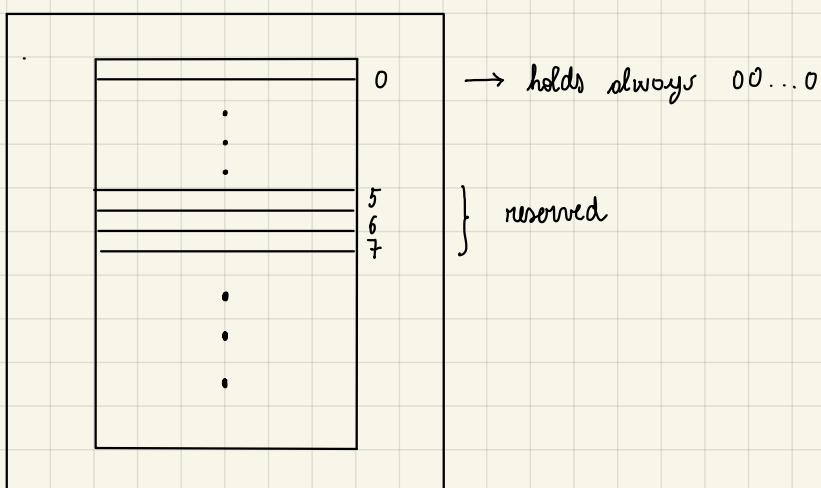
- In RISC-V every instruction is stored in one word



- Fetch and execute

25/02/2022

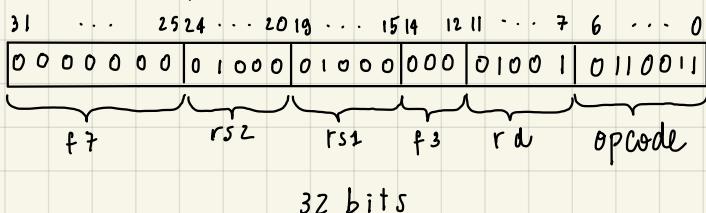
REGISTERS IN CPU



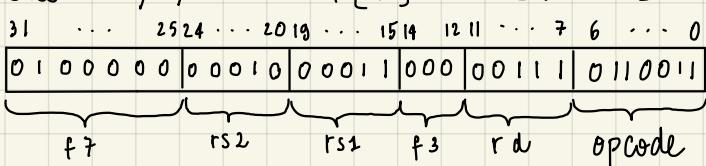
INSTRUCTIONS IN RISC-V ARCHITECTURE

word	
1	Instr. 1
2	0x00400000
3	In Str. 2
4	.
5	Instr. 3
6	.
7	Instr. 4
	.

① Add $5, 8, 8 \Rightarrow R[5] \leftarrow R[8] + R[8]$



② Sub $7, 3, 2 \Rightarrow R[7] \leftarrow R[3] - R[2]$



Ex. 1

$$R[5] \leftarrow R[28] + R[29] + R[30]$$

Add 5, 28, 29

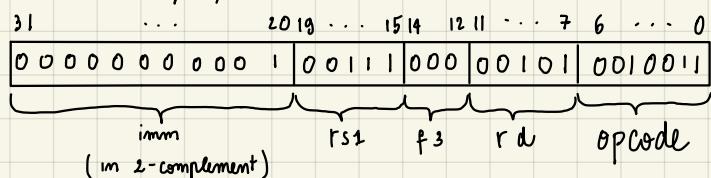
Add 5, 5, 30

Ex. 2

$$R[5] \leftarrow R[7]$$

Add 5, 7, 0

③ Addi 7, 7, 1 $\Rightarrow R[7] \leftarrow R[7] + 1$



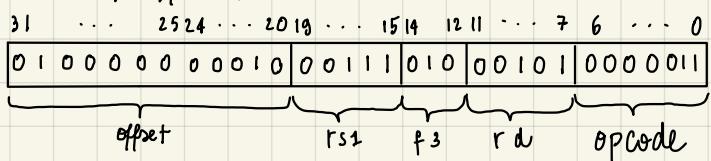
Ex.

$$R[5] \leftarrow 3$$

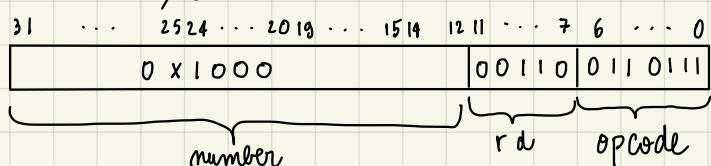
Addi 5, 0, 3

→ loads a word from the memory

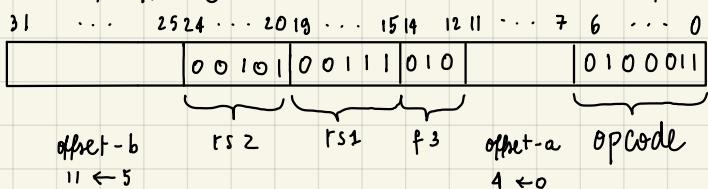
④ lw 5, offset(7) $\Rightarrow R[5] \leftarrow \text{MEM}[R[7] + \text{offset}]$



⑤ lui 6, 0x1000



⑥ $sw\ 5, \text{ offset}(8) \Rightarrow \text{MEM}[R[8] + \text{offset}] \leftarrow R[5]$



EXERCISE

$$R[5] \leftarrow x + y$$

X	0x10010000
Y	0x10010004
x+y	0x10010008

lui 1, 0x10010

lw 5, 0 (1)

lw 6, 4 (1)

Add 5, 5, 6

sw 5, 8 (1)

addi 17, 0, 10
ecall } to close the program

⑦ $beq\ 5, 6, \text{ offset} \Rightarrow \text{if } R[5] = R[6] : \text{PC} \leftarrow \text{PC} + \text{offset}$

EX.

add 5, 0, 0	0x00400000	↑
add 6, 0, 0	0x00400004	↑
beq 5, 6, 0	0x00400008	↑
addi 5, 5, 1	0x0040000C	}
addi 6, 6, 1	0x00400010	

⑧ $bne\ 5, 6, \text{ offset} \Rightarrow \text{if } R[5] \neq R[6] : \text{PC} \leftarrow \text{PC} + \text{offset}$

addi 5, 0, 0 bne 6, 0, -8 ~~loop~~

addi 6, 0, 10

loop: add 5, 5, 6

addi 6, 6, -1

09/03/2022

REGISTERS IN THE CPU

0 : x_0 zero registers

5 → 7 : $x_5 \rightarrow x_7$ $t_0 \rightarrow t_2$ they usually store temporary values

10-11 : $x_{10} - x_{11}$ $a_0 - a_1$

17 : x_{17} , a_7

SYSTEM CALLS

ecall

$a_7 = 10 \Rightarrow \text{exit}()$

$a_7 = 1 \Rightarrow \text{print_int}(a_0)$

$a_7 = 4 \Rightarrow \text{print_string}(a_0)$
containing an address in memory where the string starts

LOGICAL INSTRUCTIONS

and $x_7, x_5, x_6 \Rightarrow x_7 \leftarrow x_5 \text{ AND } x_6$

andi $x_7, x_5, 0x01 \Rightarrow x_7 \leftarrow x_5 \text{ AND } 0x01$

or ..

ori ..

xor ..

xori

11/03/2022

ble / blt $t_0, t_1, \text{label} \Rightarrow \text{PC goes to label if } t_0 < \leq t_1$

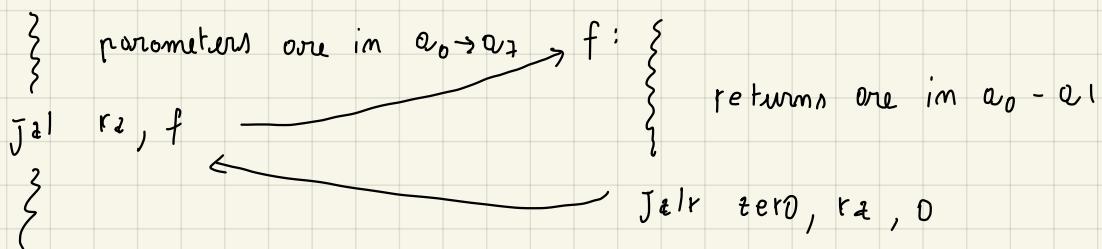
bge / bgt $t_0, t_1, \text{label} \Rightarrow \text{PC goes to label if } t_0 > \geq t_1$

FUNCTIONS

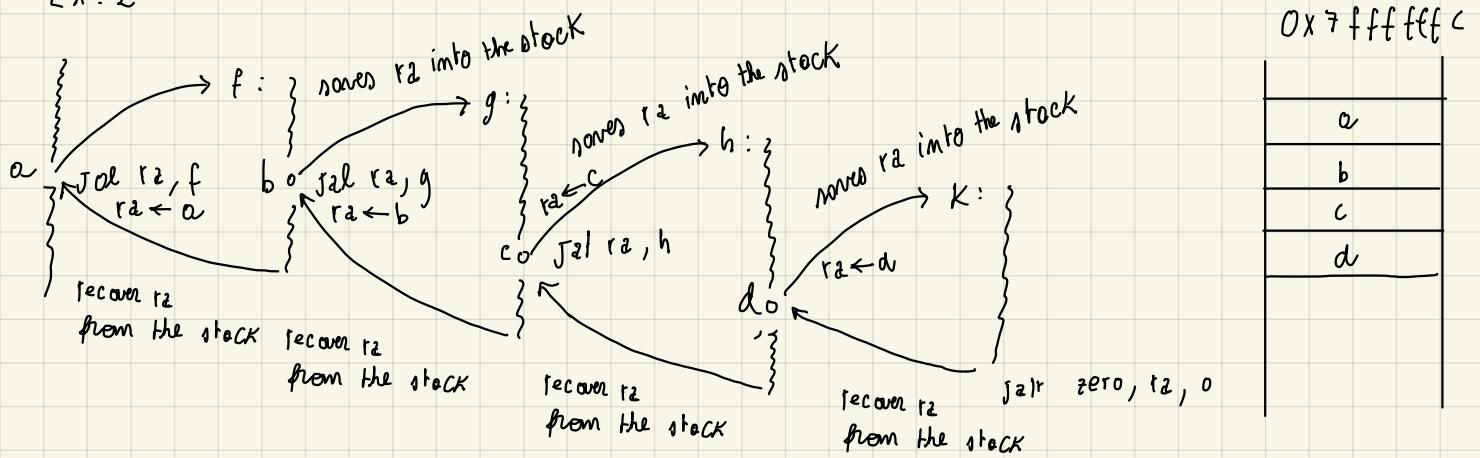
jal $r_2, \text{offset} \Rightarrow r_2 \leftarrow \text{PC} + 4$
 $\text{PC} \leftarrow \text{PC} + \text{offset}$

$\text{jalr } x_1, r_2, \text{ offset} \Rightarrow x_1 \leftarrow \text{PC} + 4$
 $\text{PC} \leftarrow r_2 + \text{offset}$

Ex. 1



Ex. 2



IMPLEMENTATION OF A STACK

Ex. 1 \leftarrow stack pointer

addi SP, SP, -4

sw r2, 0(SP)

lw r2, 0(SP)

addi SP, SP, +4

Ex. 2

addi SP, SP, -8

sw r2, 0(SP)

sw r0, 4(SP)

lw r2, 0(SP)

lw r0, 4(SP)

addi SP, SP, 8

FACTORIAL

f: bne a0, zero, rec

li 20, 1

jztr zero, r2, 0

rec: addi sp, sp, -8

sw r2, 0(sp)

sw a0, 4(sp)

addi a0, a0, -1

jz1 r2, f

20

lw t0, 4(sp)

24

mul a0, a0, t0

28

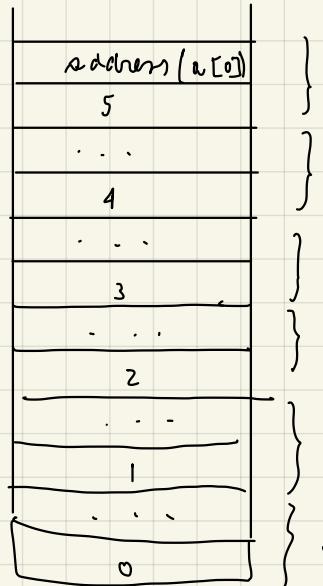
lw r2, 0(sp)

32

addi sp, sp, 8

jztr zero, r2, 0

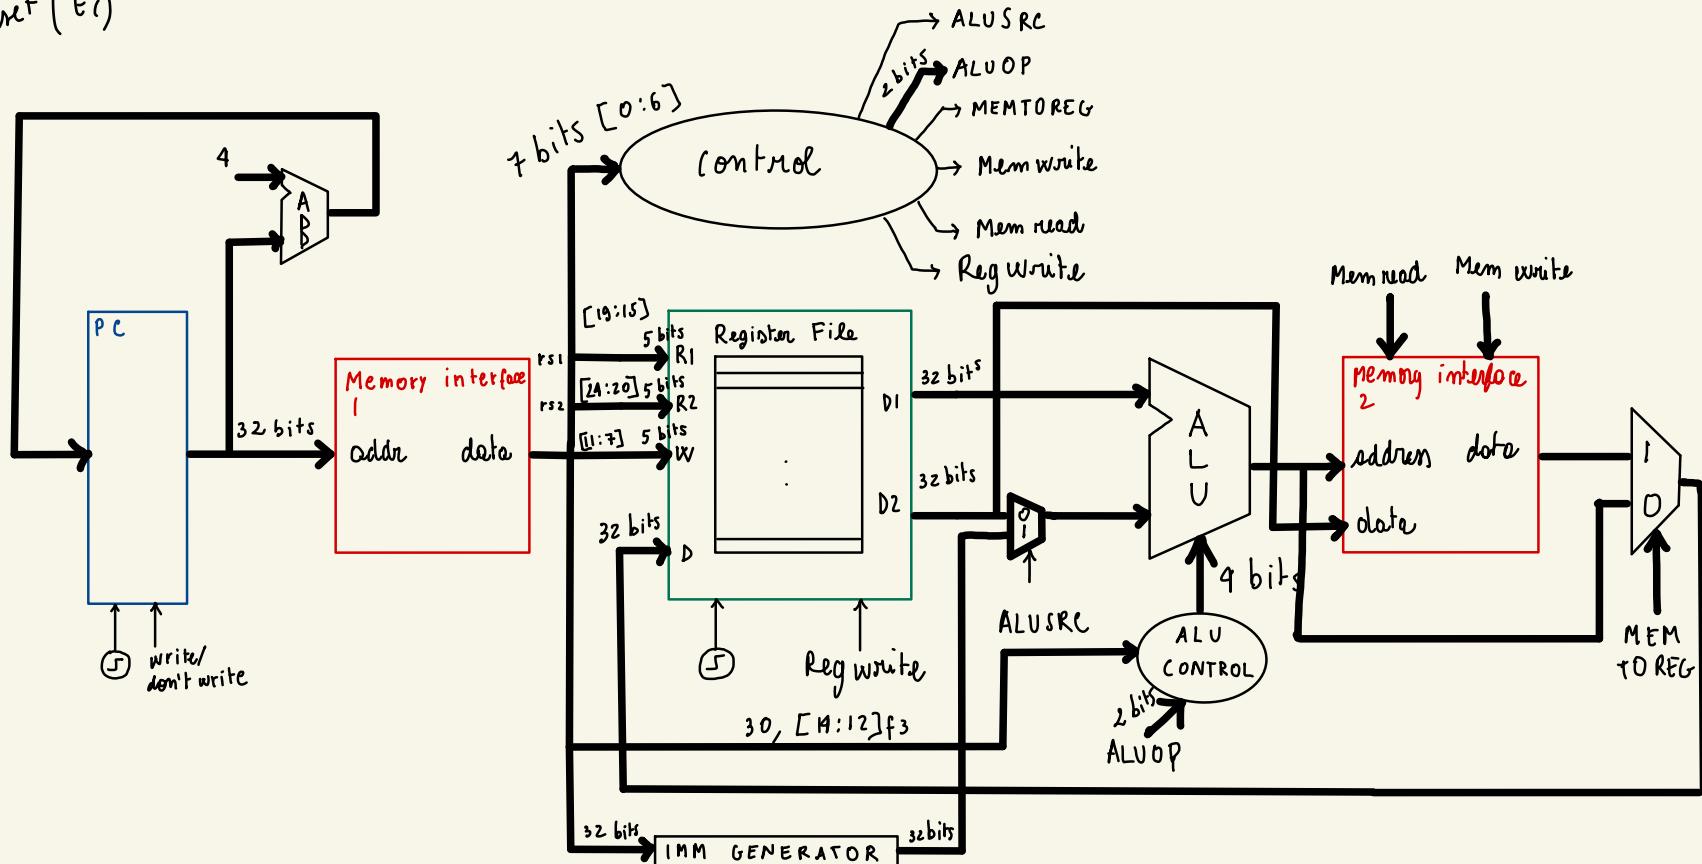
t0



16/03/2022

add	t_0, t_1, t_2
sub	t_0, t_1, t_2
addi	t_0, t_1, imm
lw	$t_0, \text{offset}(t_1)$
sw	$t_0, \text{offset}(t_1)$

SINGLE CYCLE ARCHITECTURE

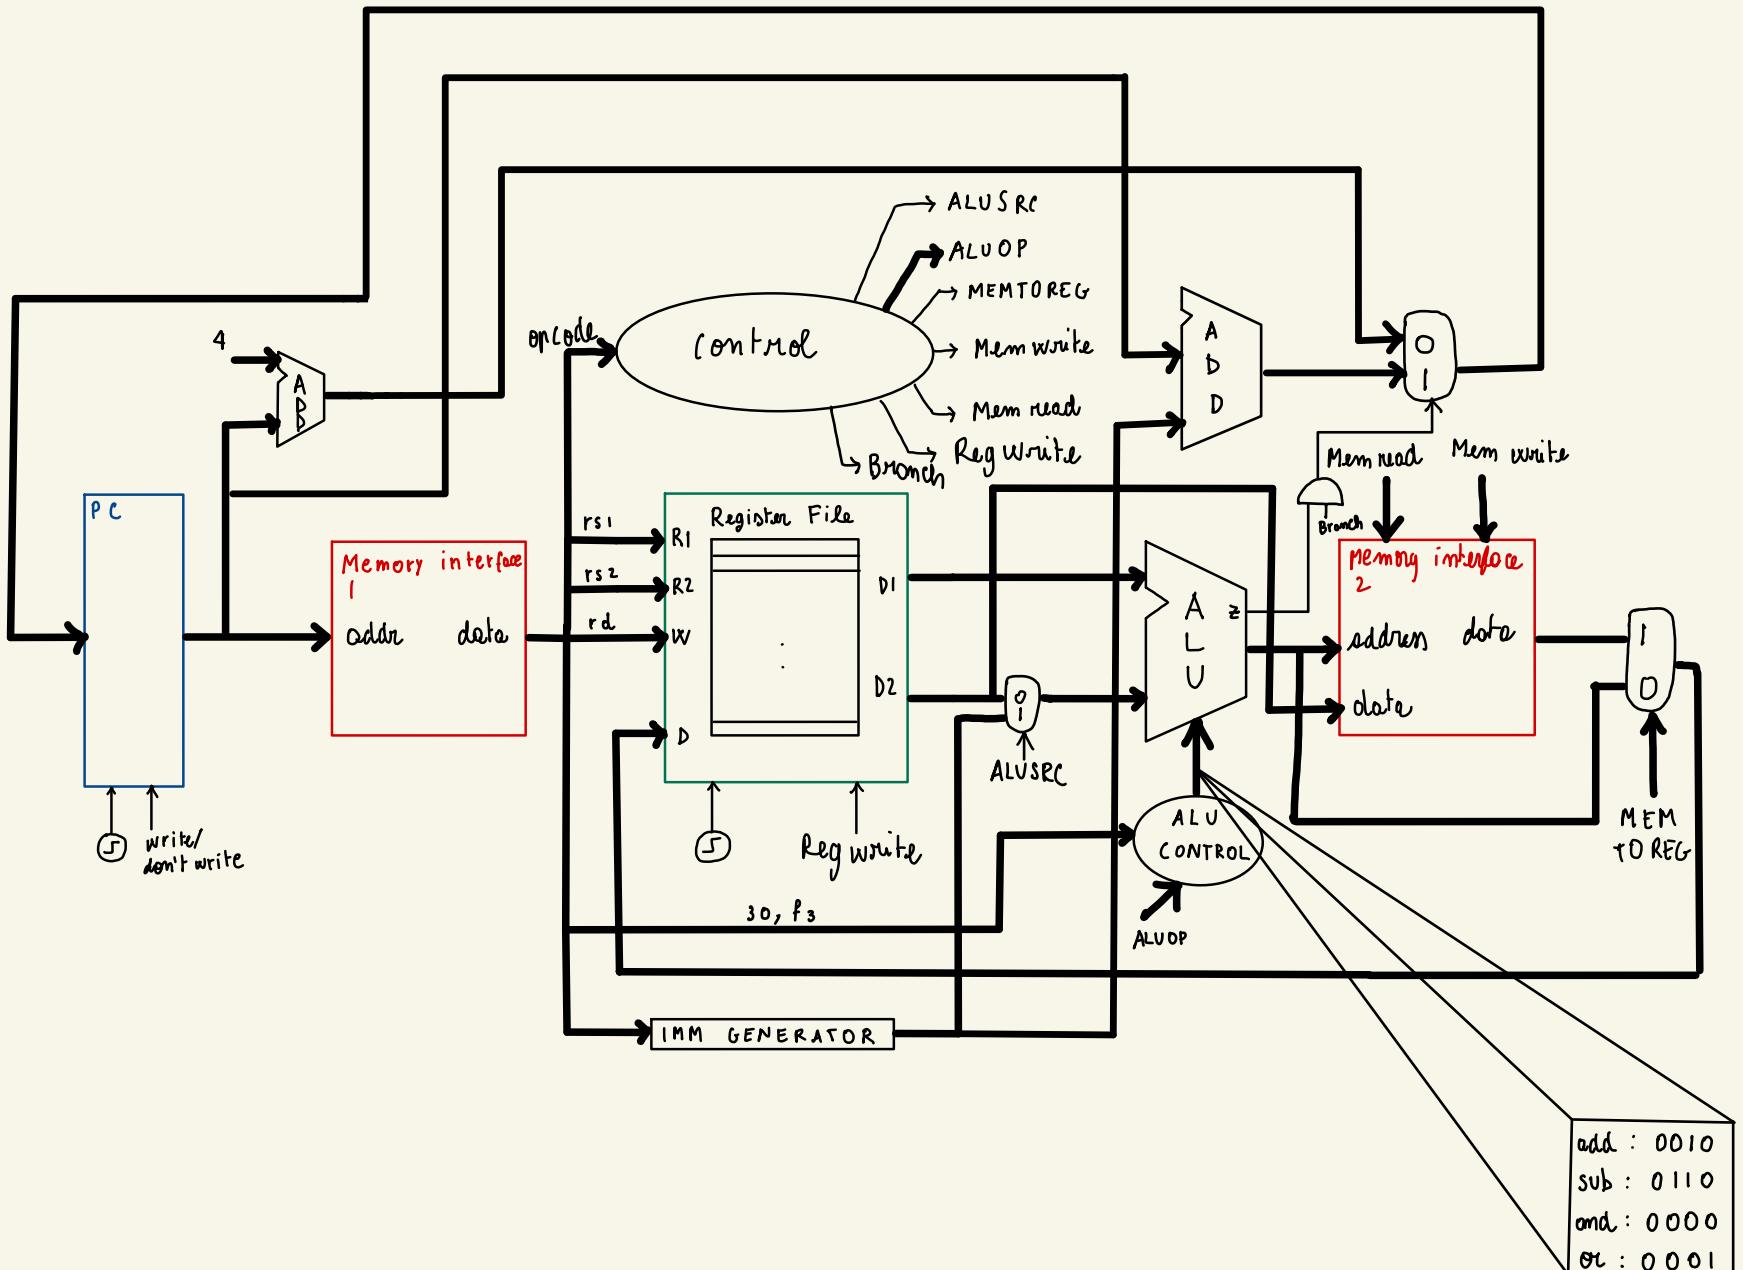


23/03/2021

add t_0, t_1, t_2
 sub t_0, t_1, t_2
 addi t_0, t_1, imm
 lw $t_0, \text{offset}(t_1)$
 sw $t_0, \text{offset}(t_1)$
 beq t_1, t_2, addr

beq x_7, x_8, offset

SINGLE CYCLE ARCHITECTURE



Ex 1

$t_0 : 0x10010000$

$t_1 : 0xF F F F F F F F$

$t_2 : 0x0 0 0 0 0 0 0 0 5$

$t_3 : 0xF F F F F F F F$

Ex. 2

lw $t_0, x = lui t_0, 0x10010$

addi $t_0, t_0, 24$

data

· ascii "Alonsodro"

· align 2

· word 5, 7, 8

x: · word 6

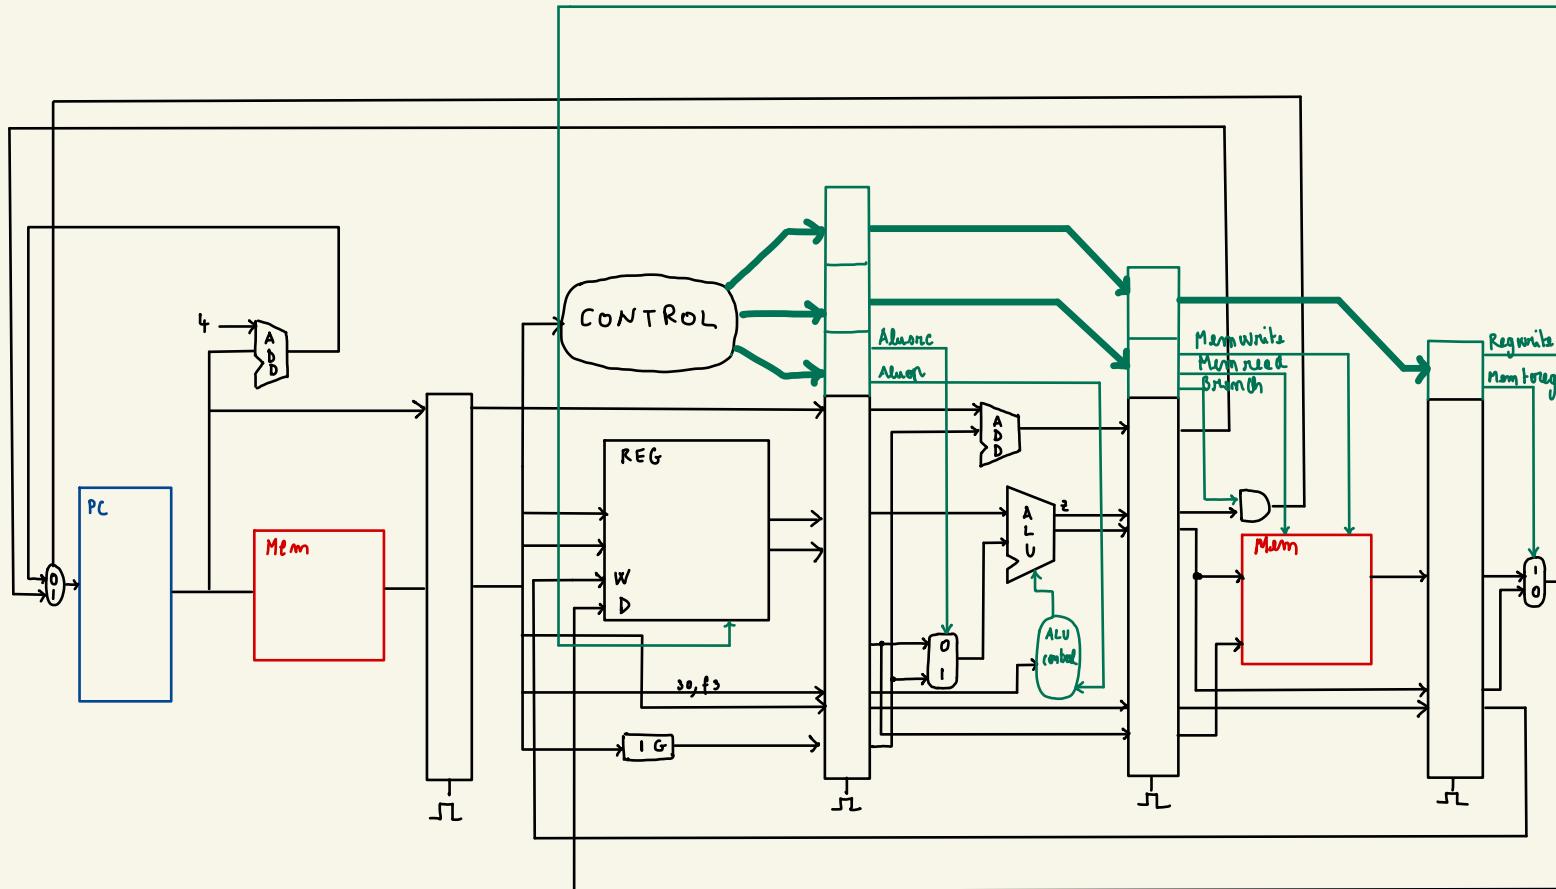
lw $t_0, x = lui t_0, 0x10010$

lw $t_0, 24(t_0)$

13/04/2022

PARALLELISM

PIPELINING



DATA HAZARDS

EX. 1

PROBLEM

add x_2, x_3, x_5

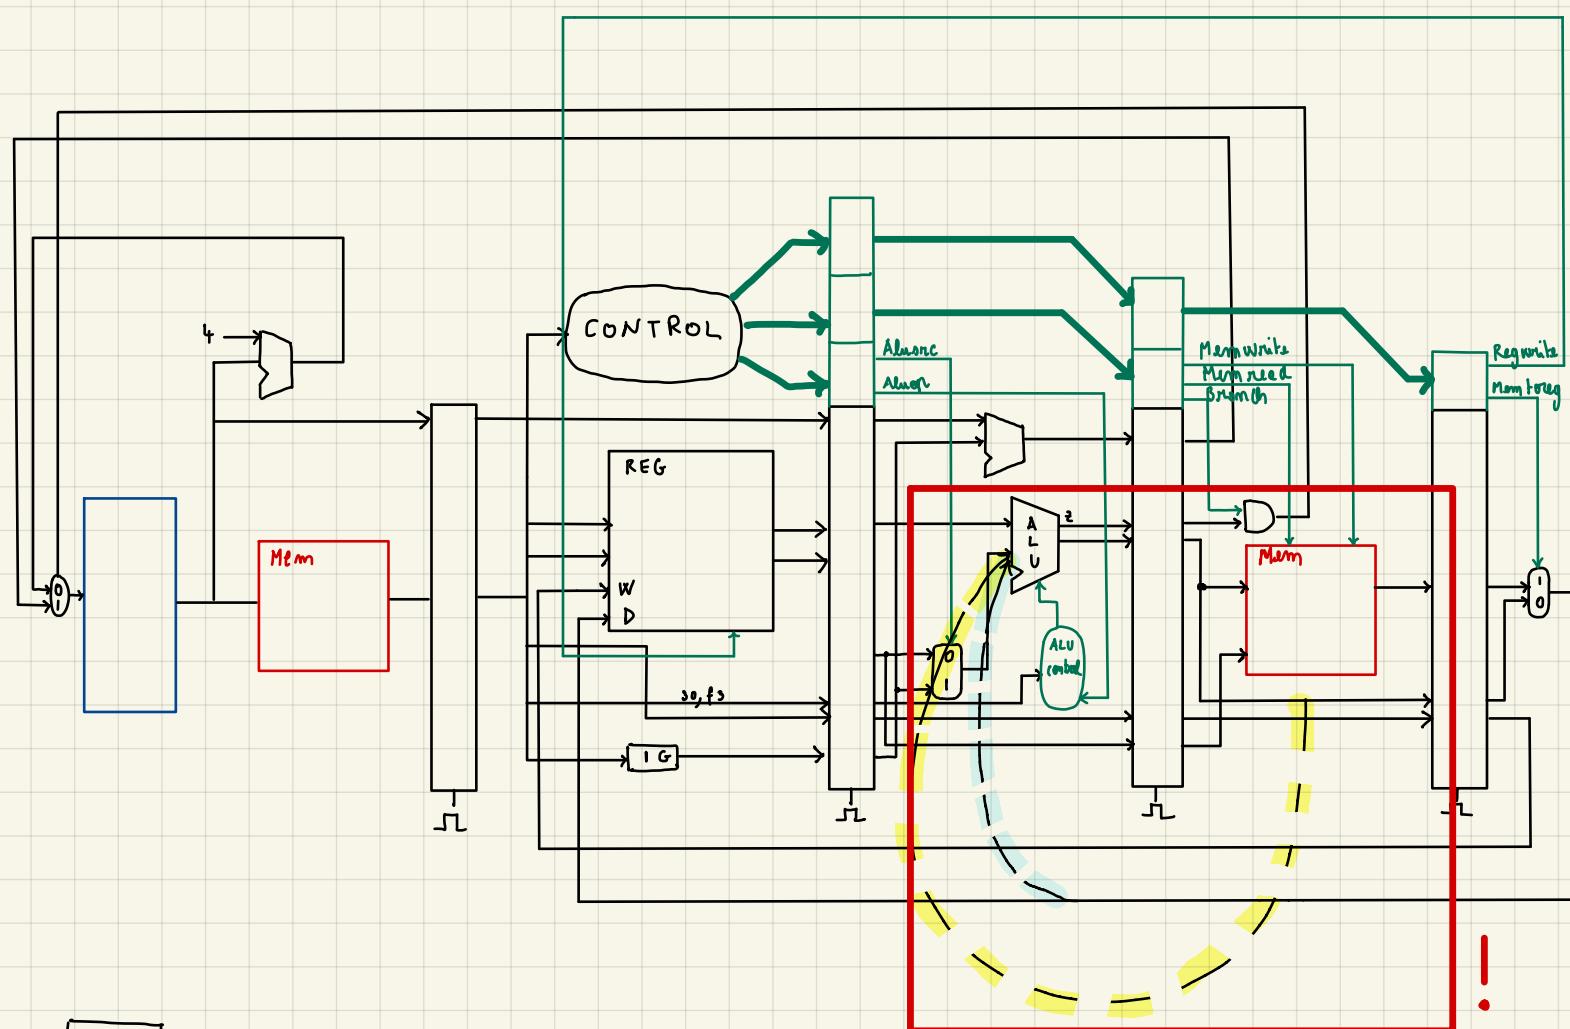
the second instruction needs the first one to complete first!

1st solution (less efficient)

Put some "bubbles" between the two instruction in order to let the first instruction complete.

2nd solution (Forwarding)

Enable the architecture to detect when a situation of that type happens and put the "future" value of x_2 in the place of the "old one"



EX. 2

PROBLEM

beq x_1, x_2, label

The architecture will continue to execute the instructions after the branch even if it should not

1st solution (loop unrolling)

reduce the number of branches

Ex.

for i in range (1, 100) → (compiler)
do-stuff()
for i in range (1, 100) :
do-stuff
do-stuff
:
do-stuff

2nd solution

Execute both the next instruction and the instruction at label address and discard the wrong one.

3rd solution

Predict before whether the branch will happen or not.

20 / 04 / 2022

FORWARDING

Ex. 1

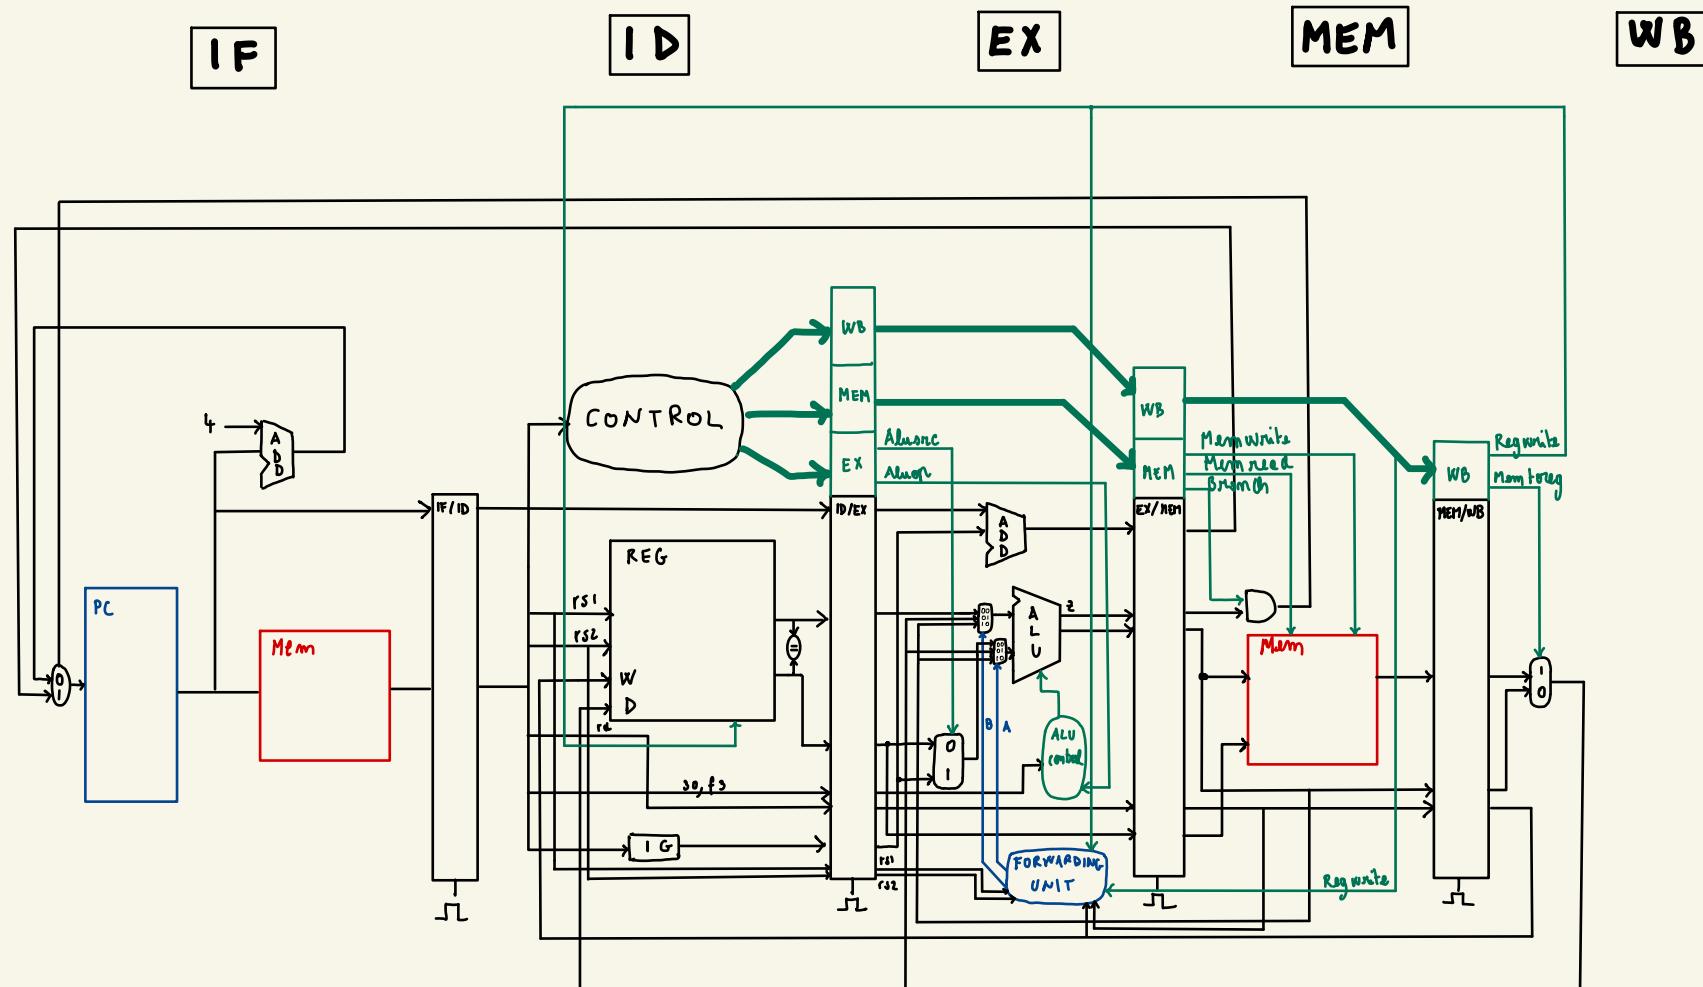
add x_1, x_1, x_2
add x_2, x_2, x_1

Ex. 2

add x_1, x_1, x_2
addi $x_7, x_7, 1$
sub x_2, x_1, x_4

Ex. 3

add x_1, x_1, x_2
add x_1, x_1, x_3
add x_1, x_1, x_4



Forwarding Unit logic

if (EX/MEM . Reg write) and (EX.MEM . rd ≠ 0) and (ID/EX.rs1 = EX.MEM . rd) :

Forward A = 10

if (EX/MEM . Reg write) and (EX.MEM . rd ≠ 0) and (ID/EX.rs2 = EX.MEM . rd) :

Forward B = 10

if (MEM/WB . Reg write) and (MEM/WB . rd ≠ 0) and (ID/EX.rs1 = MEM/WB . rd) and

not((EX/MEM . Reg write) and (EX.MEM . rd ≠ 0) and (ID/EX.rs1 = EX.MEM . rd)) :

Forward A = 01

if (MEM/WB . Reg write) and (MEM/WB . rd ≠ 0) and (ID/EX.rs2 = MEM/WB . rd) and

not((EX/MEM . Reg write) and (EX.MEM . rd ≠ 0) and (ID/EX.rs2 = EX.MEM . rd)) .

Forward B = 01

22/04/2022

STALLING

Ex.

lw x6, 0(x7) bubble

IF

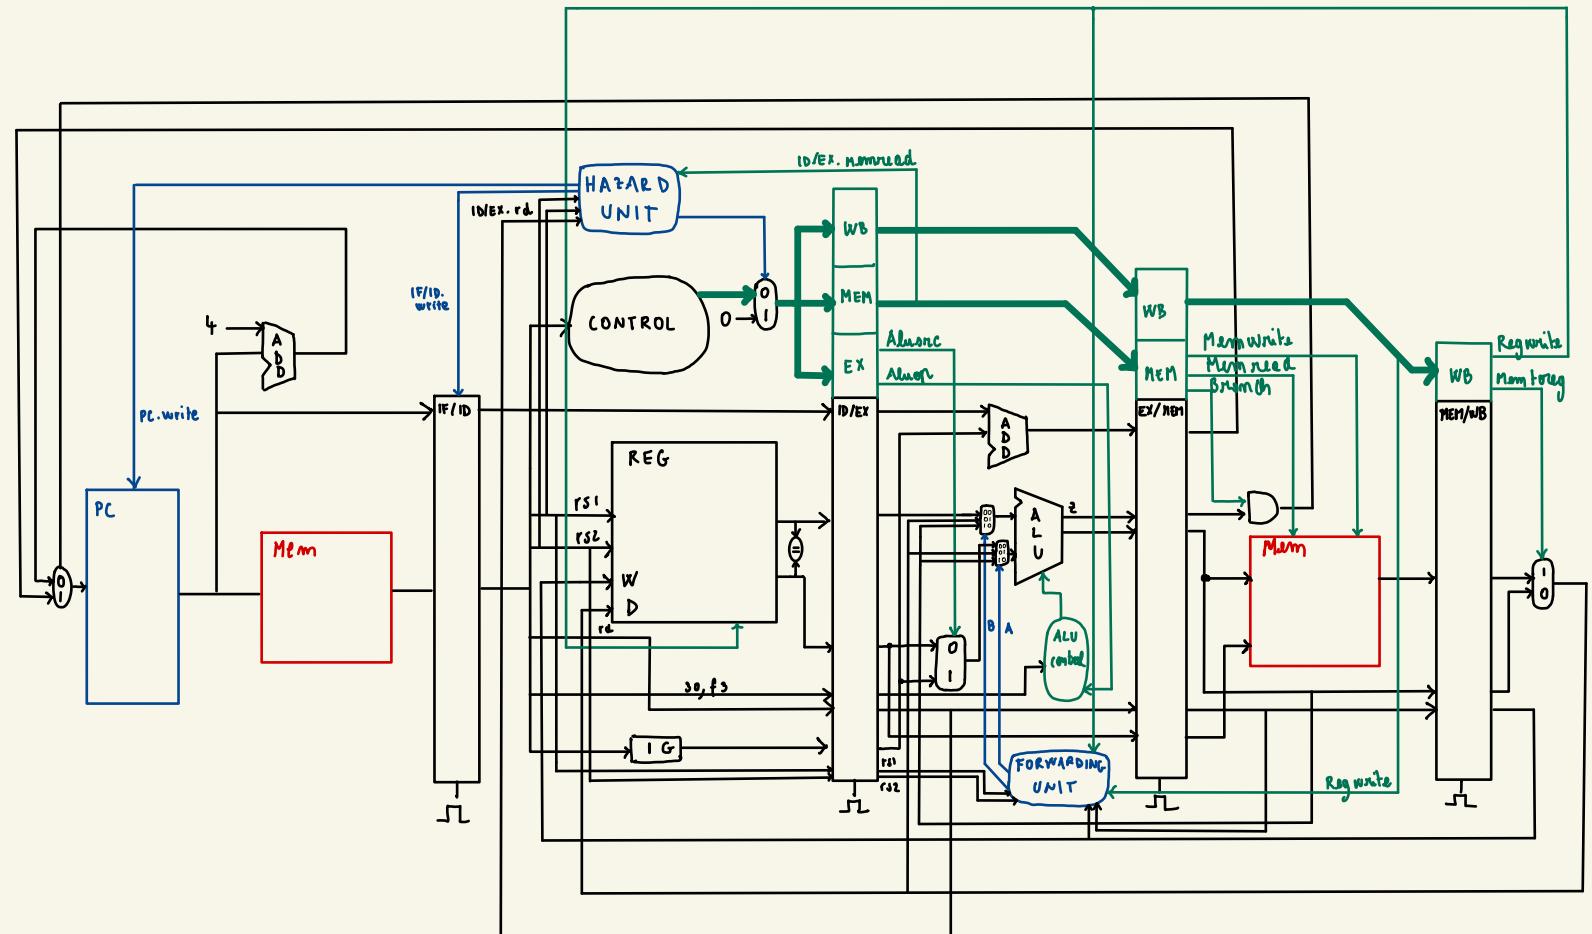
add x8, x8, x6

ID

EX

MEM

WB



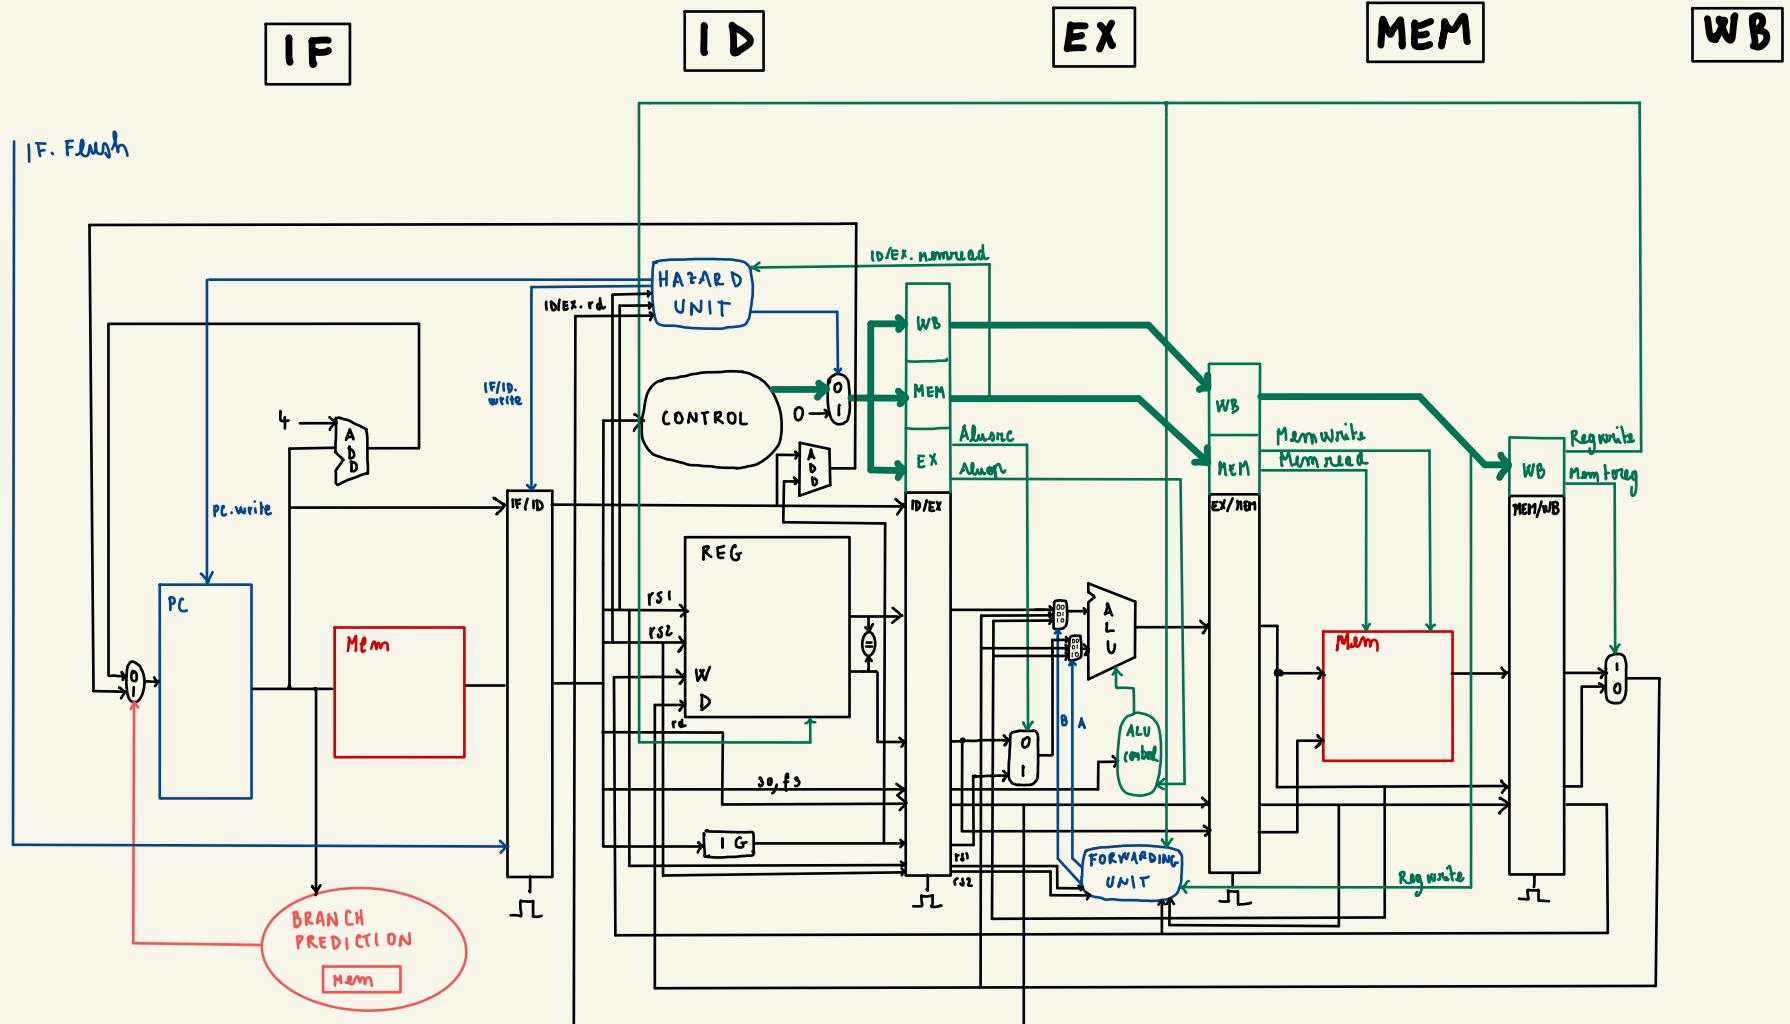
HAZARD UNIT LOGIC

if $ID/EX.\text{Memread} \text{ and } (ID/EX.\text{Md} = IF/ID.\text{RS1} \text{ or } ID/EX.\text{Md} = IF/ID.\text{RS2})$:
stall the pipeline

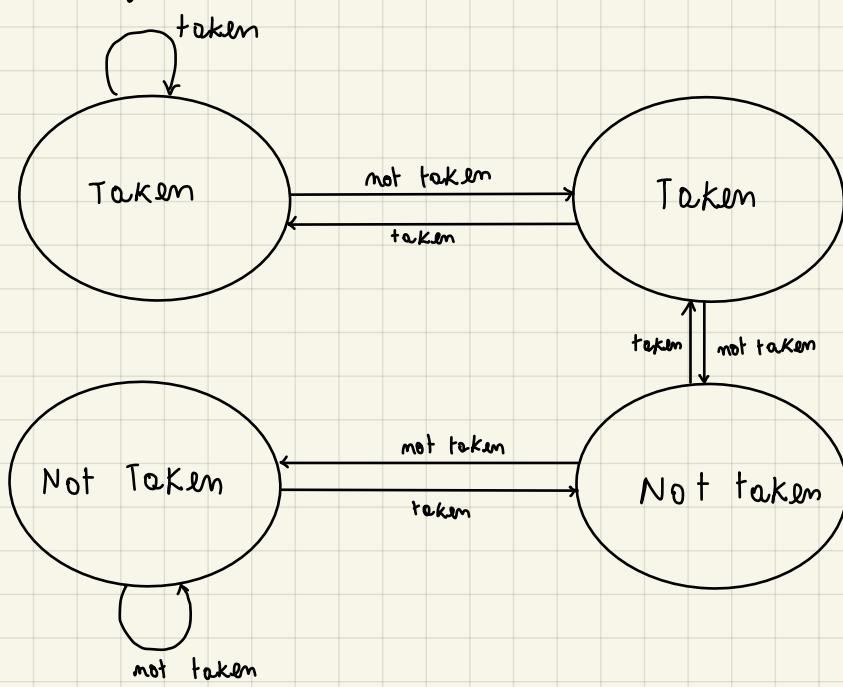
BRANCH PREDICTION

EX.

for i=1 to 1000 000



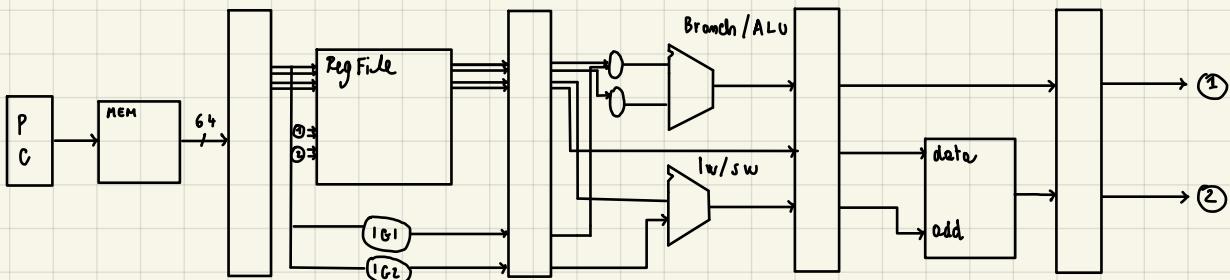
Branch prediction logic



MULTIPLE ISSUE ARCHITECTURE

pocket

static architecture



EXERCISE

.data

x: (array 1)
y: (array 2)
z: (array 3)
h: een

.text

le t₀, x
le t₁, y
le t₂, z
lw a₁, h

loop : lw s₀, 0(t₀)
lw s₁, 0(t₁)

add s₁, s₀, s₁
sw s₁, 0(t₂)
add t₀, t₀, 4
add t₁, t₁, 4
add t₂, t₂, 4
add a₁, a₁, -1
bne a₁, zero, loop

Multiple issue version (loop only)

addi t₀, t₀, 4 | lw s₀, 0(t₀)
addi t₁, t₁, 4 | lw s₁, 0(t₁)
addi a₁, a₁, -1 | /
add s₁, s₀, s₁ | /
addi t₂, t₂, 4 | sw s₁, 0(t₂)
bne a₁, zero, loop | /

Multiple issue version loop unrolling (loop only)

addi a₁, a₁, -2 | lw s₀, 0(t₀)
addi t₂, t₂, 8 | lw s₁, 0(t₁)
addi t₀, t₀, 8 | lw s₂, 4(t₀)
add t₃, s₁, s₂ | lw s₃, 4(t₁)
add t₁, t₁, 8 | sw t₃, -8(t₂)
add t₄, s₂, s₃
bne a₁, zero, loop | sw t₄, -4(t₂)

Performance for array of 10^6 items

① Performance of single clock cycle 200 MHz

$$\frac{\text{clocks per loop}}{200 \cdot 10^6} = \frac{g \cdot 10^6}{200 \cdot 10^9 \frac{1}{s}} = 0,045 \text{ s} = 45 \text{ ms}$$

② Pipeline 1 GHz
no branch prediction

$$\frac{11 \cdot 10^6}{1 \cdot 10^9 \frac{1}{s}} = 11 \cdot 10^{-3} \text{ s} = 11 \text{ ms}$$

③ Pipeline 1 GHz
branch prediction
code reordering

$$\frac{g \cdot 10^6}{1 \cdot 10^9 \frac{1}{s}} = 5 \text{ ms}$$

④ Multiple issue no branch prediction 1 GHz

$$6 \cdot 10^6 \text{ ms} = 6 \text{ ms}$$

⑤ Multiple issue no branch prediction loop unrolling 1 GHz

$$6 \cdot \frac{10^6}{2} \text{ ms} = 3,5 \text{ ms}$$

27/04/2022

MEMORY

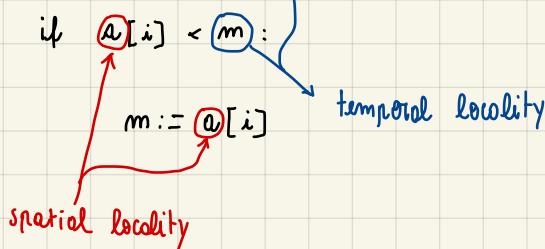
LOCALITY

- spatial locality
- temporal locality

EX.

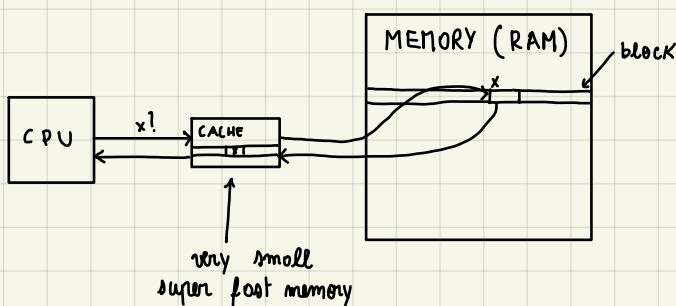
$m := a[0]$

for $i := 1$ to 1000 000 :



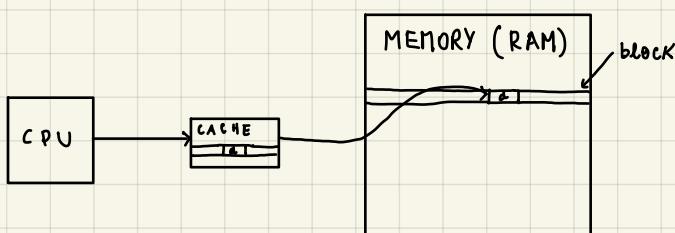
CACHE MEMORY

Read



- (i) The value asked by the CPU is already in the cache \rightarrow hit \rightarrow faster access
- (ii) Else \rightarrow miss \rightarrow the value is saved in the cache (together with his block) for future access

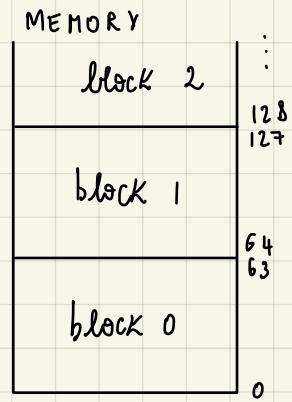
Write



- (i) The value is saved in the cache
- (ii) If the cache needs to be overwritten, the value is then saved in the memory

1 block is 64 bytes

There are 2^k sets in the cache

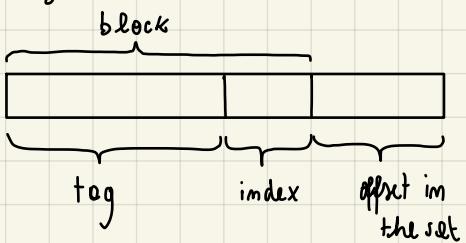


$x = \text{address}$

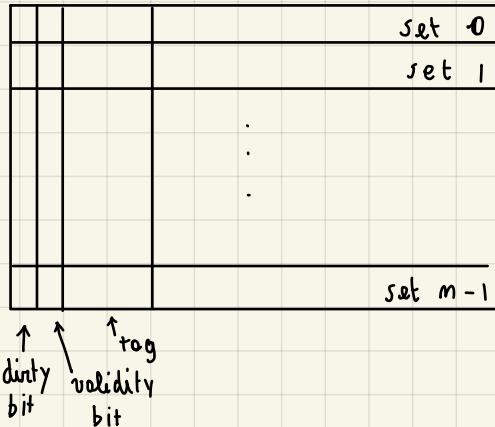
$$\text{block} = [x / 64] = x[6:]$$

$$\text{set} = \text{block \% msets} = \text{block}[:4]$$

$$\text{tag} = [\text{block} / \text{msets}]$$



CACHE



EXERCISE 4

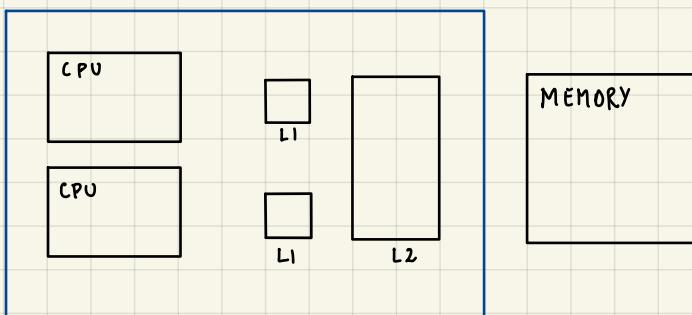
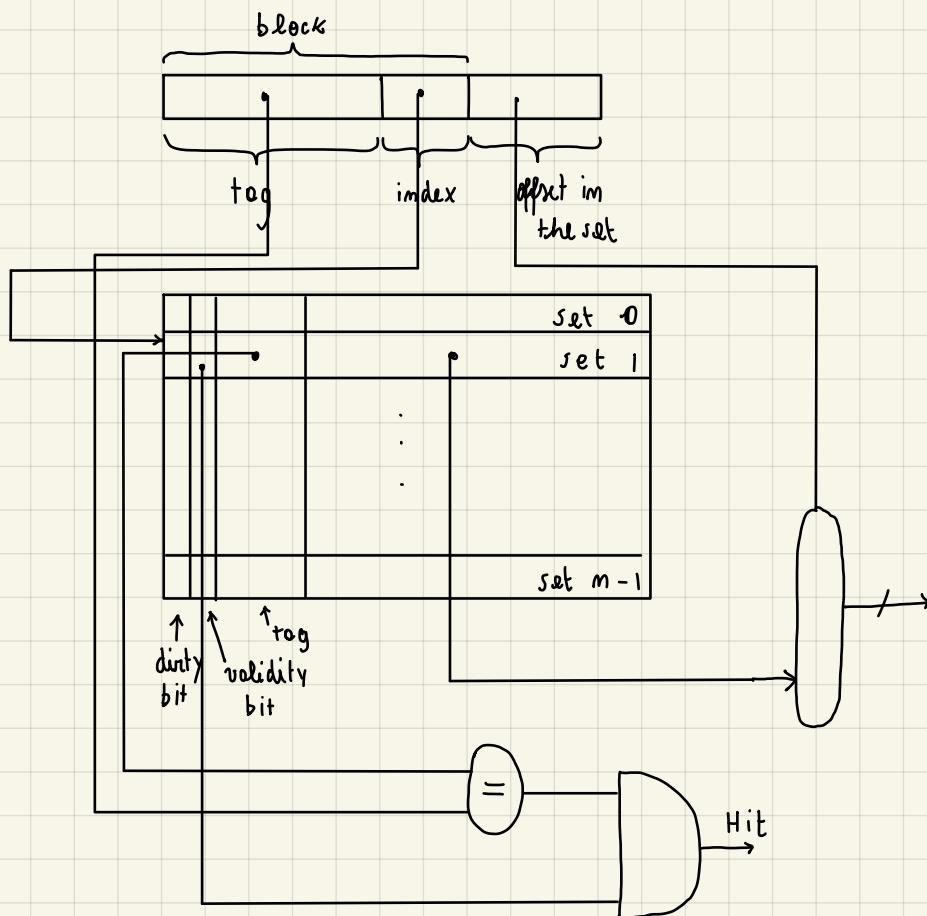
- cache with 4 cache
- block size of 32 bytes

access chronology : 36, 48, 0, 100, 164, 172, 96, 92
1 2 3 4 5 6 7 8
M H M M M H H

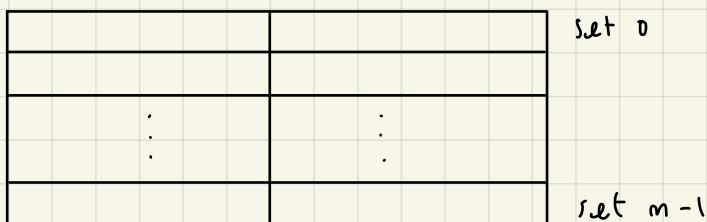
- 1) Miss \rightarrow block 1 goes in set 1 of the cache
- 2) Hit (block 1 already in the cache)
- 3) Miss \rightarrow block 0 goes into set 0
- 4) Miss \rightarrow block 3 goes into set 3
- 5) Miss \rightarrow block 5 goes into set 1
- 6) Hit (block 5 already in the cache)
- 7) Hit (block 3 already in the cache)
- 8) Miss \rightarrow block 2 goes into set 2

29/04/2022

CACHE



2-WAY ASSOCIATIVE CACHE



Ex.

.text

{

add : lw t₂, 0(t₀)

add t₃, t₃, t₂

addi t₀, t₀, 4

addi t₁, t₁, -1

bne t₁, zero, l1c

- block size = 64 bytes

- two-way associative cache with 16 sets

- LRU (least recently used) : the block which was used the least recently is discarded

data 16

Analysis (two-way cache)

ACCESES MISS

iteration 0 6 1

iteration 1 6 0

:

:

iteration 16 6 1

iteration 17 6 0

:

iteration 31 6 0

iteration 32 6 1

cycle	data 0
data 1	
data 2	
...	
...	
.	
.	
...	

Miss rate in the first 16 iterations = $\frac{1}{16 \cdot 6} \approx 1\%$
(it is the same for the others)

Analysis (one-way cache)

ACCESES MISS

iteration 0 6 2

iteration 1 6 2

:

iteration 15 6 2

iteration 16 6 1

iteration 17	6	0
:	:	:
iteration 31	6	0
iteration 32	:	

$$\text{miss rate} = \frac{32 + 31(1)}{16 \cdot 6 + 31(16 \cdot 6)} = 2,4\%$$

04/05/2022

data

x: . word	_____	0x10010000	set 0
y: . word	_____		set 0
z: . word	_____		set 0
m: . word	4096		

text

M	lw	t ₀ ,	x	0x00400000	0x00010000	block : 64
H	lw	t ₁ ,	y			set : 32
H	lw	t ₂ ,	z			0
H	lw	t ₃ ,	m			one-way association write-back

cico : M lw t₄, o(t₀) M
M lw t₅, o(t₁) M
M add t₆, t₄, t₅
H sw t₆, o(t₂) M
M addi t₀, t₀, 4
H / t₁ t₁ /
H / t₂ t₂ /
H / t₃ t₃ -1
M bne t₃, zero, cico } → set 1

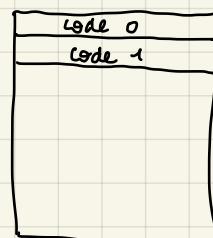
i 0	8	12
i 1	6	12
:		
i 15	6	12
i 16	4	12
:		
i 31	4	12
i 32	3	12
i 33	3	12
:		

i2) H M M M M N M M L H H H H

i16) H M H M H H H M H H H H

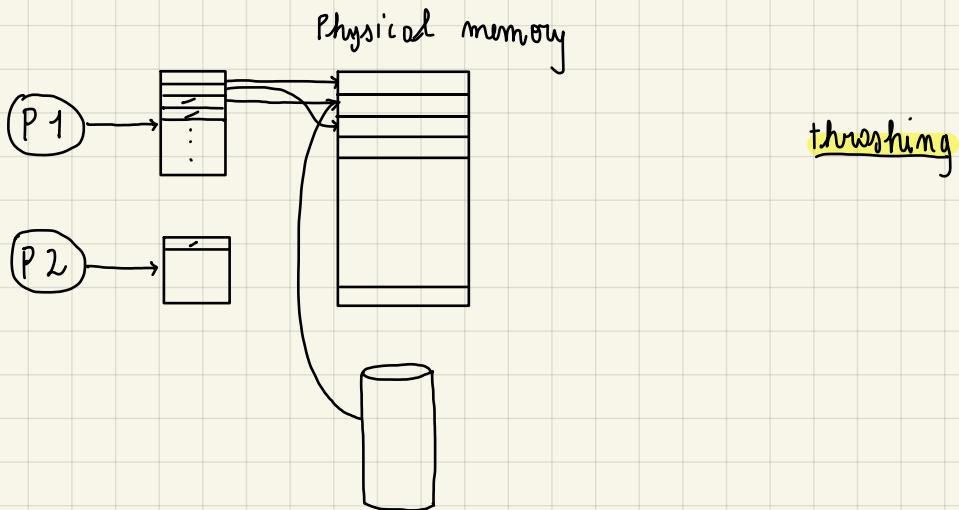
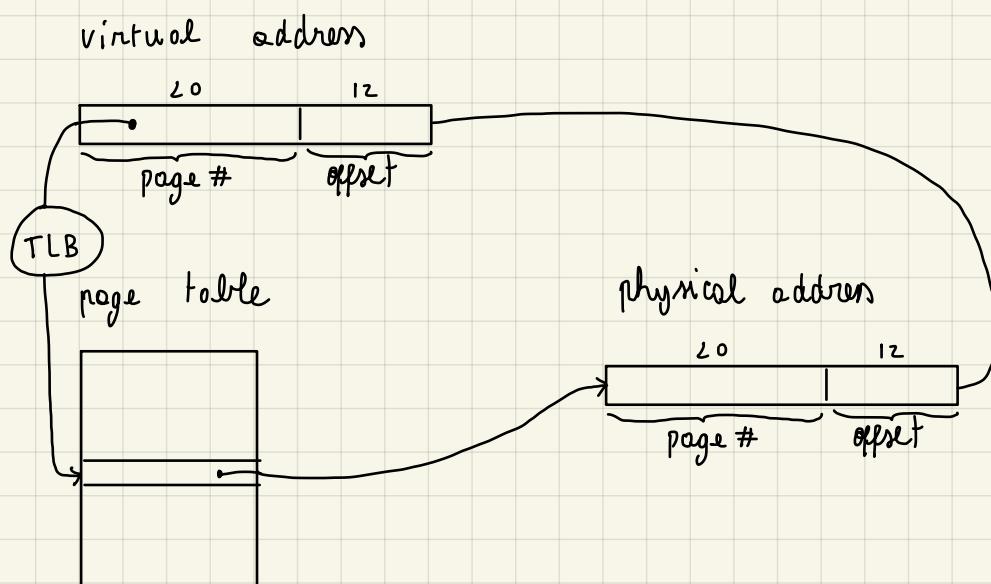
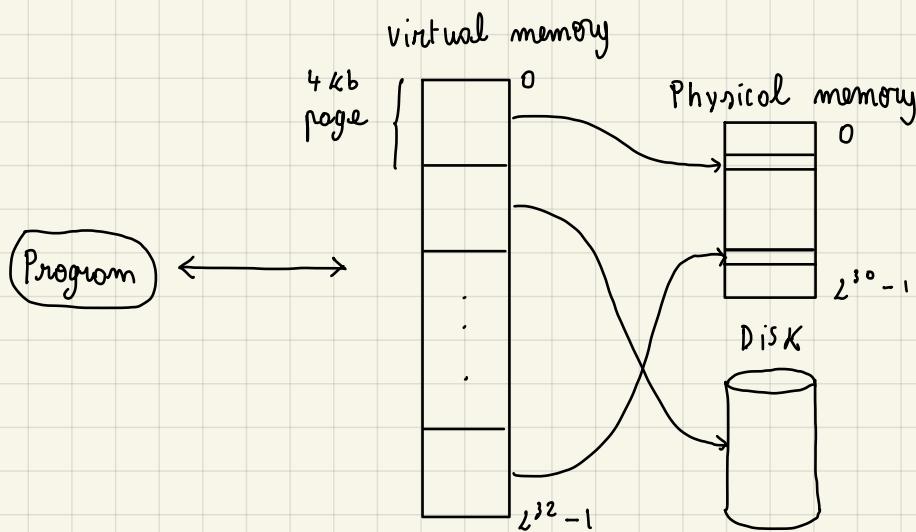
i32) H M H M N H M H H H H

$$\text{miss rate} = \frac{8 + 15 \cdot 6 + 4 \cdot 16 + 3 \cdot 16 \cdot 30}{12 \cdot 16 + 12 \cdot 16 + 12 \cdot 16 \cdot 30} = 2,6\%$$



11/05/2022

VIRTUAL MEMORY

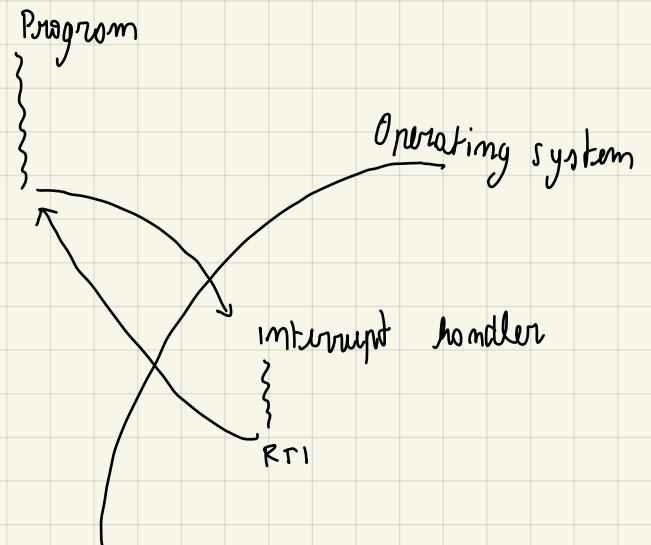


13/05/2022

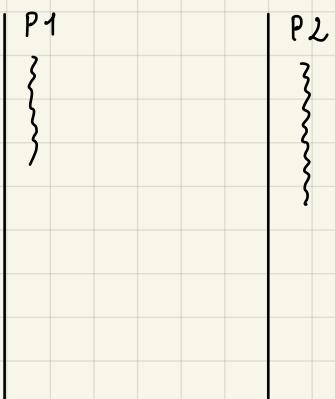
EXCEPTIONS / INTERRUPTS

EXAMPLES

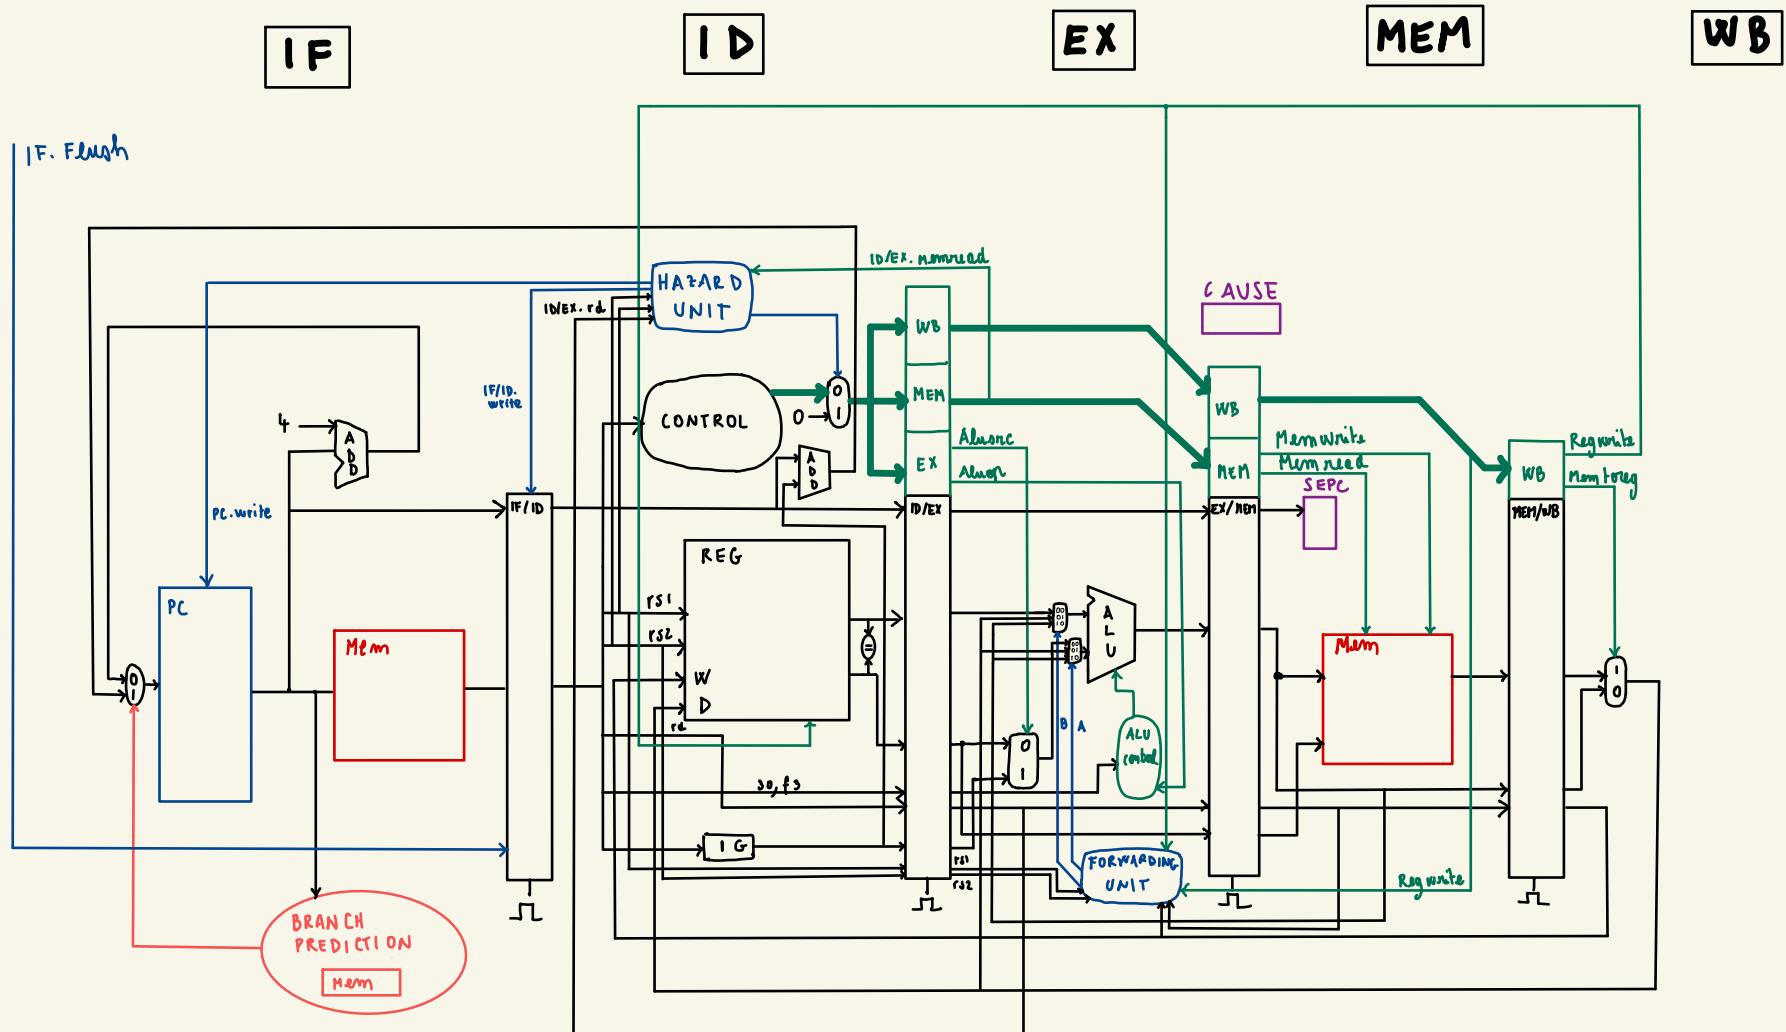
- signal to the process
- division by zero
- overflow
- call
- undefined instruction
- memory violation
- hardware fault
- end of time slice
- I/O operations



TIME SLICE



vector interrupt



I/O

DATA FROM I/O DEVICES TO CPU

Two ways:

- 1) I/O instructions (not very used)
- 2) Memory mapped I/O

Ex.

$\begin{matrix} 0x00000000 \\ 0x00010000 \end{matrix}$ } addresses reserved to the devices

INTERACTION BETWEEN I/O DEVICES AND

Polling : CPU asks periodically to the device if it has completed

interrupt driven I/O : I/O device sends a signal to the CPU after it has completed

PARALLELISM

SISD (single instruction single data)

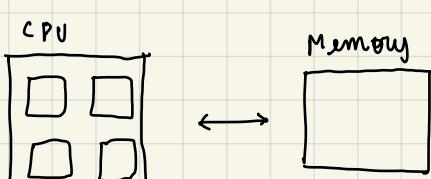
SIMD (- / multiple /)

MISD (multiple / single /)

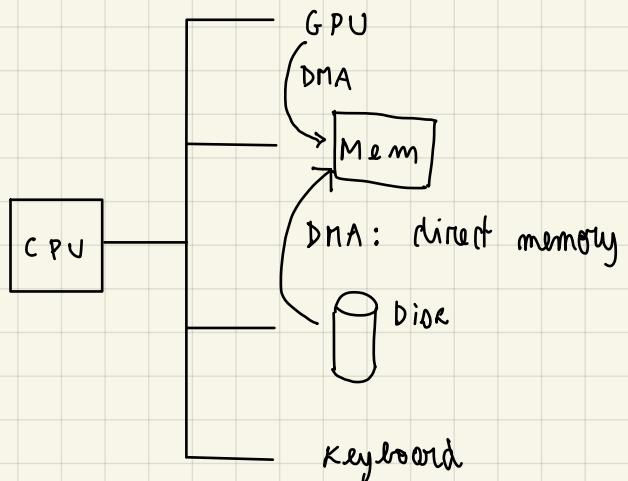
MIMD (multiple / multiple /) → multicore

MIMD

Multiple CPUs
Multiple cores

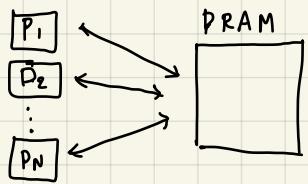


Multithreading



SIMD

Theoretical PRAM



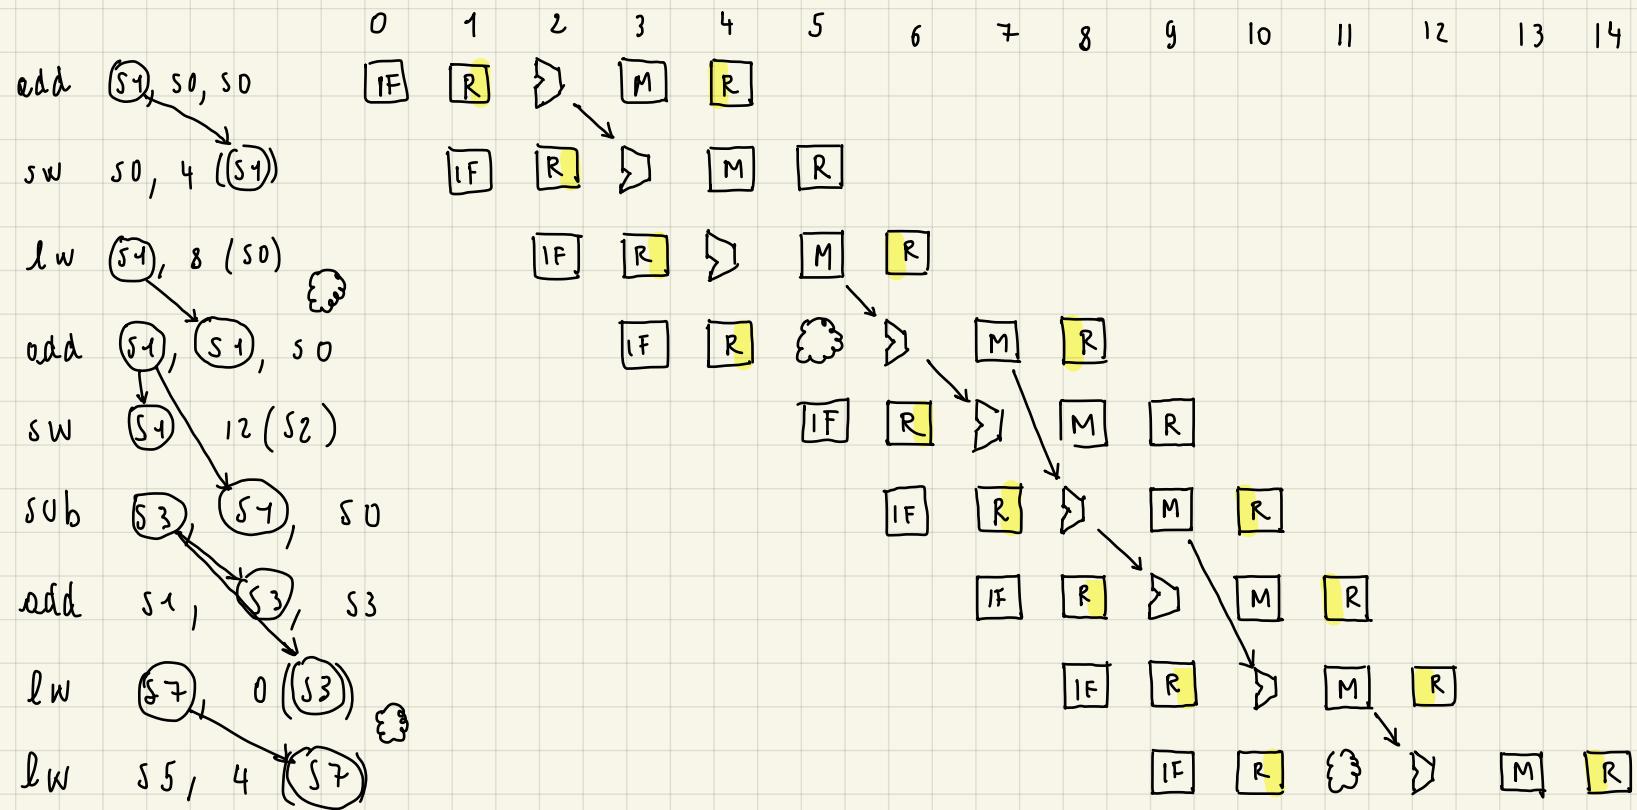
20/05/2022

Exercise 1. How many clock cycles are needed to complete the following code on the Risc-V pipelined architecture? (Figure in the next page.) Remember that:

- Execution completes when the last instruction exits the pipeline.
- If you read and write the same register of the register file in the same clock cycle, the value that is read is the value that is being written.

Show a figure of the pipeline with bubbles and forwardings.

```
add $1, $0, $0
sw $0, 4($1)
lw $1, 8($0)
add $1, $1, $0
sw $1, 12($2)
sub $3, $1, $0
add $1, $3, $3
lw $7, 0($3)
lw $5, 4($7)
```



60

Exercise 2. Consider a two-way associative cache with four sets and block of 8 words. The cache is initially empty and replacement of blocks is done by using LRU. Show which of the following memory accesses are hit and which are misses: 128, 16, 4, 184, 68, 204, 104, 28, 252, 240, 244, 40, 136, 164, 56, 256, 152, 180, 156, 64, 48.

block = 32 bytes

address	block	set	H/M
128	4	0	M

0	4	8
1	5	1
2	2	6
3	3	7

16	0	0	M
----	---	---	---

0 6

48	0	0	H
----	---	---	---

1 32

68	2	2	M
----	---	---	---

2 64

204	6	2	M
-----	---	---	---

3 96

104	3	3	M
-----	---	---	---

4 128

28	0	0	H
----	---	---	---

5 160

252	7	3	M
-----	---	---	---

6 192

240	7	3	H
-----	---	---	---

244	7	3	H
-----	---	---	---

40	1	1	M
----	---	---	---

136	4	0	H
-----	---	---	---

164	5	1	H
-----	---	---	---

56	1	1	H
----	---	---	---

256	8	0	M
-----	---	---	---

152	4	0	H
-----	---	---	---

180	5	1	H
-----	---	---	---

156	4	0	H
-----	---	---	---

64	2	2	H
----	---	---	---

48	1	1	H
----	---	---	---

Exercise 3. Consider the following code:

```
.data
v:    .word 0,1,7,4,2,...      # array of 1024 integers  data 0
w:    .word 0,0,0,0,0,...      # array of 1024 integers  data 64
n:    .word 1024

.text
    lui s0, 0x10012          code 0
    lw s2, 0(s0)
    lui s0, 0x10010
    lui s1, 0x10011
loop:   lw t0, 0(s0)
        sw t0, 0(s1)
        addi s2, s2, -1
        bne s2, zero, loop
        addi s0, s0, 4
        addi s1, s1, 4
        ecall
        li a7, 10
        ecall
```

This program copy one array into another.

- Question A:** what is the approximate total miss rate if you have an instruction cache and a data cache, both one-way associative with 8 blocks of 64 bytes?
- Question B:** What is the speed-up (how faster is it) if we use the Risc-V multiple issue architecture with loop unrolling of 4 loops instead of the standard pipelined architecture? Show the code.

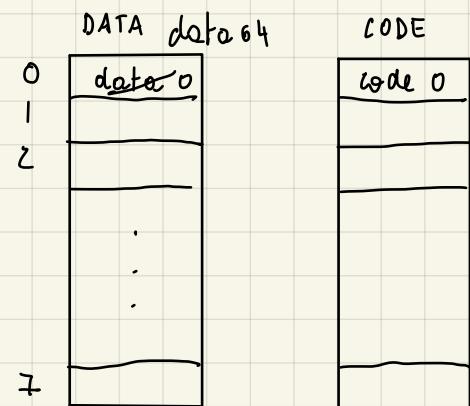
(1)

$n_0 : H M H M H H H H$

; Always the same

$n_{1023} /$

$$\text{miss rate} : \frac{4}{8} = 25\%$$



(2)

standard pipeline (branch prediction)

clock cycles (for the loop) : 6 apart first and last iteration

total clock cycles $\approx 6 \cdot 1024$

multiple issue architecture:

loop:

lw	t ₀ ,	0(\$0)	addi	s ₀ , s ₀ , 16
lw	t ₁ ,	-12(\$0)		/
lw	t ₂ ,	-8(\$0)		/
lw	t ₃ ,	-4(\$0)		/
sw	t ₀ ,	0(\$1)	addi	s ₁ , s ₁ , 16
sw	t ₁ ,	-12(\$1)		/
sw	t ₂ ,	-8(\$1)	addi	s ₂ , s ₂ , -4
sw	t ₃ ,	-4(\$1)	bne	s ₂ , zero, loop

$$\text{total clock cycles} \approx 8 \cdot \frac{1024}{4} = 2 \cdot 1024$$

$$\frac{6 \cdot 1024}{2 \cdot 1024} = 3 \Rightarrow \text{M1 is 3x faster}$$

.data

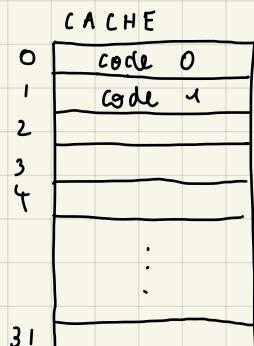
x: . word	—	0x10010000	data 0	block = $[x / 64]$
y: . word	—	0x10014000	data 256	
z: . word	—	0x10018000	data 512	set = block % m sets
m: . word	$4096 \cdot 4 = 2^{14}$			

.text

l0: t₀, x 0x00400000
 l0: t₁, y
 l0: t₂, z
 lw t₃, m

cicle:
 lw t₄, o(t₀)
 lw t₅, o(t₁)
 add t₆, t₄, t₅
 sw t₆, o(t₂)
 addi t₀, t₀, 4
 / t₁ t₁ /
 / t₂ t₂ /
 / t₃ t₃ -1
 bne t₃, zero, cicle

(16 words)
 block size : 64 bytes
 number of sets : 32
 one-way associative
 write-back



memory access per cycle: 12

MISSES

i0: M M M M M H M M M H H H M 8

i9: H M M M M H M M H H H H 6

:

i15

x8 i16 H M H M H H M H H H H 4

:

,

i32 H M H M H H M H H H H 3

:

i511

i512

$$\frac{8 + 6 \cdot 15 + 4 \cdot 16 + 3 \cdot 480}{512 \cdot 12} = 26\%$$

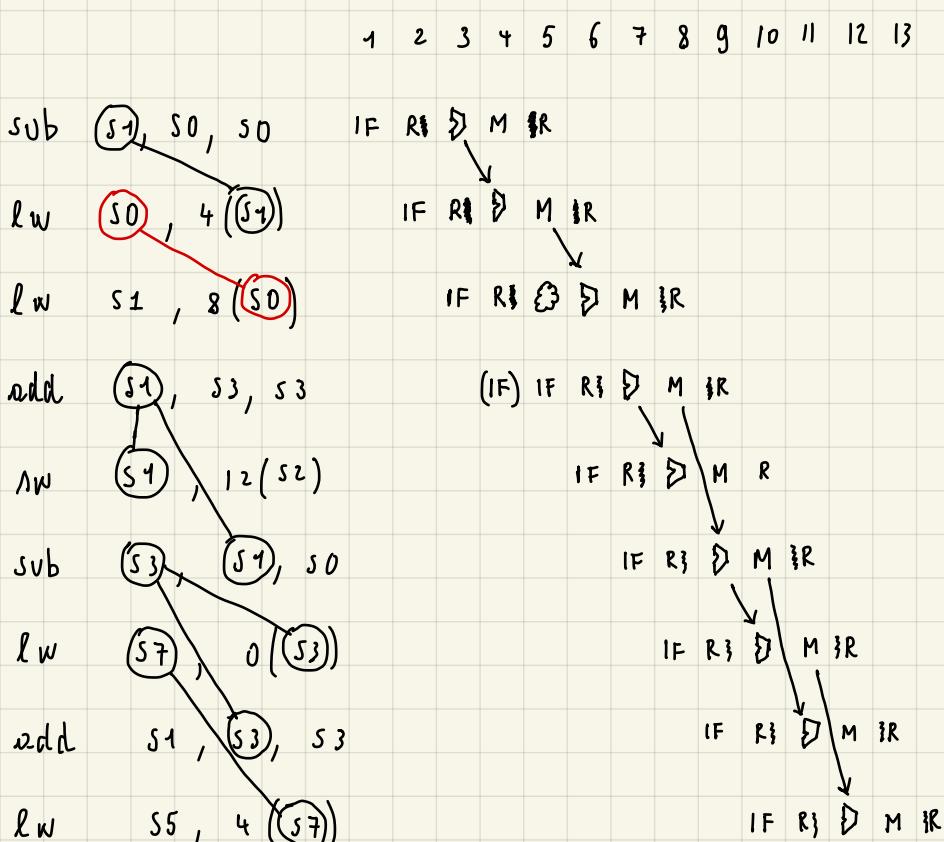
SAMPLE EXAM 2

Exercise 1. How many clock cycles are needed to complete the following code on the Risc-V pipelined architecture? (Figure in the next page.) Remember that:

- Execution completes when the last instruction exits the pipeline.
- If you read and write the same register of the register file in the same clock cycle, the value that is read is the value that is being written.

Show a figure of the pipeline with bubbles and forwardings.

```
sub s1, s0, s0
lw s0, 4(s1)
lw s1, 8(s0)
add s1, s3, s3
sw s1, 12(s2)
sub s3, s1, s0
lw s7, 0(s3)
add s1, s3, s3
lw s5, 4(s7)
```



Exercise 2. Consider a two-way associative cache with two sets and block of 8 words. The cache is initially empty and replacement of blocks is done by using LRU. Show which of the following memory accesses are hit and which are misses: 16, 136, 184, 252, 160, 80, 252, 80, 48, 136, 60, 204, 100, 88, 196, 60, 32, 176, 80, 224.

block size: 32 bytes

BLOCK	SET	H/M
16	0	M
136	4	M
184	5	M
252	7	M
160	5	H
80	2	M
252	7	H
80	2	H
48	1	M
136	4	H
60	1	H
204	6	M
100	3	M
88	2	M
196	6	H
60	1	H
32	1	H
176	5	M
80	2	H
224	7	M

0	6	2
1	1	5

0	0
32	1
64	2
96	3
128	4
160	5
192	6
224	7
256	8

Exercise 3. Consider the following code:

```
.data
n000: .word 8, n001      # list of 1024 integers      data  0
n001: .word -3, n002
n002: .word 5, n003
[...]

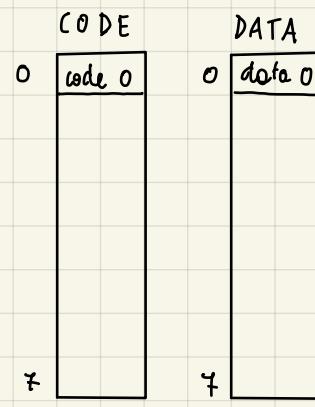
lista: .word n000

.text
    lui s0, 0x10012
    lw s1, 0(s0)
    li s2, 0
loop:   beq s1, zero, fine
        lw s1, 4(s1)
        addi s2, s2, 1
        bne s1, zero, loop
fine:   ecall
        li a7, 10
        ecall
```

This program counts the number of nodes in the list.

- Question A:** what is the approximate total miss rate if you have an instruction cache and a data cache, both one-way associative with 8 blocks of 64 bytes?
- Question B:** What is the speed-up (how faster is it) if we use the Risc-V multiple issue architecture with loop unrolling of 4 loops instead of the standard pipelined architecture? Show the code.

①
i 0 : H H M H H
i 1 : H H H H H
:
i 8 : H H H H H



$$\text{miss rate} = \frac{1}{40} = 2,5\%$$

②

```
beq s1, zero, fine | lw s1, 4(s1)
addi s2, s2, 1 |
beq s1, zero, fine | lw s1, 4(s1)
addi s2, s2, 1 |
beq s1, zero, fine | lw s1, 4(s1)
```

addi \$2, \$2, 1 | -

bne \$4, zero, fine | lw \$4, 4(\$4)

addi \$2, \$2, 1 | -

bne \$4, zero, loop | -

$$\frac{4 \cdot 1024}{9 \cdot 256} = \frac{16 \cdot 256}{9 \cdot 256} = 1, \overline{7}$$

M1 is 1,7x faster